

Homework 6: Transductive Support Vector Machines

Christian Bock
christian.bock@bsse.ethz.ch

Dr. Michael Moor
michael.moor@bsse.ethz.ch

Leslie O'Bray
leslie.obray@bsse.ethz.ch

Prof. Dr. Karsten Borgwardt
karsten.borgwardt@bsse.ethz.ch

Submission deadline: 02.06.2021 at 14.00

Objectives

The goals of this homework are:

- to implement a variant of SVM using the **cvxopt** package.
- to implement a transductive SVM [Joachims, 1999] with label switching strategy.
- to compare inductive and transductive SVM [Joachims, 1999] on a text dataset.

Dataset

You will work on a text dataset extracted from Reuters articles compiled by Thorsten Joachims ¹. All documents are represented as feature vectors. Each feature corresponds to the tf-idf values of a word (9930 features). The dataset consists of only 10 training examples (5 positive and 5 negative) and 600 test examples. The dataset is stored in two files **train.dat** and **test.dat**. You will find the 600 test documents in **test.dat**, where the first column indicates the true class label (either 1 or -1). You will find the 10 training documents and the same 600 test documents in **train.dat**, where the first column indicates the class label (1, -1 or 0, where 0 means unlabeled test documents). The dataset is stored in sparse format. The script for reading the data is given in **run_hw6.py**.

Package

In order to run the code, the following Python dependencies have to be fulfilled: **numpy**, **scikitlearn**, **cvxopt**, where **cvxopt** is a python package for convex optimization. You will find the information of **cvxopt** in the following link.

<http://cvxopt.org>

¹http://download.joachims.org/svm_light/examples/example2.tar.gz

Exercise 1

The primal problem of soft margin SVM [Cortes and Vapnik, 1995] is:

$$\begin{aligned} \min_{w,b,\xi} \quad & \frac{1}{2} \|\vec{w}\|^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t.} \quad & y_i(\langle \vec{w}, \vec{x}_i \rangle + b) \geq 1 - \xi_i, \quad \forall i \in \{1, \dots, n\} \\ & \xi_i \geq 0, \quad \forall i \in \{1, \dots, n\} \end{aligned}$$

Its dual representation is:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \langle \vec{x}_i, \vec{x}_j \rangle \\ \text{s.t.} \quad & \sum_{i=1}^n \alpha_i y_i = 0, \\ & 0 \leq \alpha_i \leq C, \quad \forall i \in \{1, \dots, n\} \end{aligned}$$

The dual representation is a quadratic programming problem:

$$\begin{aligned} \min_{\alpha} \quad & \frac{1}{2} \alpha^T P \alpha + q^T \alpha \\ \text{s.t.} \quad & G \alpha \leq h, \\ & A \alpha = b \end{aligned}$$

where

$$P = y \otimes y \cdot X X^T, q = -\vec{1}, A = y^T, b = 0, G = \begin{bmatrix} -I \\ I \end{bmatrix} h = \begin{bmatrix} \vec{0} \\ C \cdot \vec{1} \end{bmatrix}$$

Given P, q, A, b, G, h , soft margin SVM can be solved using `cvxopt` package. The implementation is given in the function `train()` of file `svm_linear.py`.

Exercise 1.a The variant of SVM used in transductive SVM [Joachims, 1999] is:

$$\begin{aligned} \min_{w,b,\xi,\xi^*} \quad & \frac{1}{2} \|\vec{w}\|^2 + C \sum_{i=1}^n \xi_i + C_-^* \sum_{j:y_j^*=-1} \xi_j^* + C_+^* \sum_{j:y_j^*=1} \xi_j^* \\ \text{s.t.} \quad & y_i(\langle \vec{w}, \vec{x}_i \rangle + b) \geq 1 - \xi_i, \quad \forall i \in \{1, \dots, n\} \\ & y_j^*(\langle \vec{w}, \vec{x}_j \rangle + b) \geq 1 - \xi_j^*, \quad \forall j \in \{1, \dots, k\} \\ & \xi_i \geq 0, \quad \forall i \in \{1, \dots, n\} \\ & \xi_j^* \geq 0, \quad \forall j \in \{1, \dots, k\} \end{aligned}$$

Prove that its dual representation is:

$$\begin{aligned}
 \max_{\alpha} W(\alpha) &= \sum_{i=1}^{n+k} \alpha_i - \frac{1}{2} \sum_{i=1}^{n+k} \sum_{j=1}^{n+k} \alpha_i \alpha_j y_i y_j \langle \vec{x}_i, \vec{x}_j \rangle \\
 \text{s.t. } \sum_{i=1}^{n+k} \alpha_i y_i &= 0, \\
 0 \leq \alpha_i &\leq C, \quad \forall i \in \{1, \dots, n\} \\
 0 \leq \alpha_i &\leq C_{-}^*, \quad \forall i \in \{n+1, \dots, n+k\} \text{ and } y_i = -1 \\
 0 \leq \alpha_i &\leq C_{+}^*, \quad \forall i \in \{n+1, \dots, n+k\} \text{ and } y_i = 1
 \end{aligned}$$

Exercise 1.b Complete the implementation of the variant of SVM [Joachims, 1999] described in **Exercise 1.a** in function `train_variant()` of the file `svm_linear.py`. Perform it and the standard soft margin SVM [Cortes and Vapnik, 1995] on the given text dataset using only labeled training documents.

Exercise 2

Exercise 2.a Implement the transductive SVM algorithm [Joachims, 1999] described in Algorithm 1 in function `train()` of the file `tsvm_linear.py`. Perform it on the given dataset using both labeled and unlabeled training documents.

Algorithm 1 Transductive SVM [Joachims, 1999]

Input: X_1 and y_1 , Labeled training samples

X_2 , Unlabeled training samples

C_1 and C_2 , Regularizers for labeled and unlabeled samples

p , Percentage of positive samples in unlabeled data X_2

Output: A weight vector \vec{w} and an intercept b

- 1: $\vec{w}, b = \text{svm.train}(X_1, y_1, C_1)$
 - 2: $\text{num}_+ = |X_2| * p$
 - 3: The num_+ samples from X_2 with the highest value of $\vec{w} * \vec{x}_j + b$ are assigned to positive class ($y_2[j] := 1$) and the others are assigned to the negative class ($y_2[j] := -1$)
 - 4: $C_2^- = 1 \times 10^{-5}$, $C_2^+ = (1 \times 10^{-5}) * \frac{\text{num}_+}{|X_2| - \text{num}_+}$
 - 5: **while** $C_2^- < C_2$ or $C_2^+ < C_2$ **do**
 - 6: $\vec{w}, b, \xi_1, \xi_2 = \text{svm.train_variant}(X_1, y_1, X_2, y_2, C_1, C_2^-, C_2^+)$
 - 7: **while** $\exists m, l : y_2[m] \times y_2[l] < 0$ and $\xi_2[l] > 0$ and $\xi_2[m] > 0$ and $\xi_2[m] + \xi_2[l] > 2$ **do**
 - 8: $y_2[m] = -1 \times y_2[m]$
 - 9: $y_2[l] = -1 \times y_2[l]$
 - 10: $\vec{w}, b, \xi_1, \xi_2 = \text{svm.train_variant}(X_1, y_1, X_2, y_2, C_1, C_2^-, C_2^+)$
 - 11: $C_2^- = \min(2C_2^-, C_2)$
 - 12: $C_2^+ = \min(2C_2^+, C_2)$
 - 13: **return** \vec{w} and b
-

Exercise 2.b We did not tune the parameters of the transductive SVM [Joachims, 1999] in this homework. How would you choose regularization parameters C_1 and C_2 ?

Exercise 2.c To make use of the kernel trick, we map our input data $x \in \mathcal{X}$ into \mathcal{H} using the map $\phi : \mathcal{X} \rightarrow \mathcal{H}$. \mathcal{H} is a space of functions called the *Reproducing Kernel Hilbert Space* (RKHS). Explain why \mathcal{H} is called *reproducing*?

Command-line arguments

You should only complete the functions in `svm_linear.py` and `tsvm_linear.py`, and not modify the file `run_hw6.py` for submission. The file `run_hw6.py` must be executable using the following command:

```
$ python run_hw6.py
```

Note: Zero points for the **programming part** of this homework if your script is **not** executable with the above command-line argument, or the running time of your script is over **10 minutes**!

Grading and submission guidelines

This homework is worth a total of 100 points. Table 1 shows the points assigned to each exercise.

50 pts.	Exercise 1
	20 pts. Exercise 1.a
	30 pts. Exercise 1.b
50 pts.	Exercise 2
	30 pts. Exercise 2.a
	10 pts. Exercise 2.b
	10 pts. Exercise 2.c

Table 1: Grading key for Homework 6

References

- C. Cortes and V. Vapnik. Support-vector networks. *Machine Learning*, 20(3):273–297, 1995. doi: 10.1007/BF00994018. URL <http://dx.doi.org/10.1007/BF00994018>.
- T. Joachims. Transductive inference for text classification using support vector machines. In *Proceedings of the Sixteenth International Conference on Machine Learning (ICML 1999)*, Bled, Slovenia, June 27 - 30, 1999, pages 200–209, 1999.